

# Announcements

- Email me the survey: See the **Announcements** page on the course web for instructions
- Today
  - Conceptual overview of distributed systems
  - Characterization of distributed systems
- Reading
  - Today: Chapter 1 of Coulouris
  - Next time: Chapter 2 of Coulouris
- Take a break around 10:15am
- Ack: Some slides are from Coulouris or Steve Ko

# Networks of computers are everywhere!

- The Internet
- Mobile phone networks
- Corporate networks
- Factory networks
- Campus networks
- Home networks
- In-car networks
- . . .

# Main motivation for distributed systems

- Sharing of resources
  - Hardware components, e.g., disks, printers, etc.
  - Software entities, e.g., files, databases, search engines, etc.
- Resources managed by servers and accessed by clients

# A distribution system is . . .

- one in which components located at networked computers communicate and coordinate their actions only by passing messages
- (another definition) a collection of entities with a common goal, each of which is autonomous, programmable, asynchronous and failure-prone, and which communicates through an unreliable communication medium
- These definitions lead to the following characteristics of distributed systems:
  - Concurrency of components
  - No global clock
  - Independent failures of components

# Concurrency

- In a network of computers, concurrent program execution is the norm, sharing resources
- The capacity of the system to handle shared resources can be increased by adding more resources to the network

# No global clock

- When programs need to cooperate, they coordinate their actions by exchanging messages
- Close coordination depends on shared time
- There are limits to the accuracy with which the computers in a network can synchronize their clocks – there is no single global notion of the correct time [Section 6.1.1 of Tanenbaum]
- This is a consequence of the fact that the only communication is by sending messages through the network

# Independent failures

- All computer systems can fail and the system designer is responsible for planning for the consequences of possible failures
- Each component of a distributed system can fail independently, leaving the others still running
- The failure may be due to a crash or a slow response

# Examples of distributed systems

- The Internet and the associated WWW
- See next slide
  - Understanding of underlying technology in these examples is central to a knowledge of modern computing




# Selected app domains with networked apps

|  |   |
|--|---|
| <i>Finance and commerce</i>                  | eCommerce (Amazon) online banking and trading                                       |
| <i>The information society</i>               | Web information and search engines, ebooks, Wikipedia; social networking (Facebook) |
| <i>Creative industries and entertainment</i> | online gaming, music and film in the home, user-generated content (YouTube)         |
| <i>Healthcare</i>                            | health informatics, on online patient records, monitoring patients                  |
| <i>Education</i>                             | e-learning, virtual learning environments, distance learning                        |
| <i>Transport and logistics</i>               | GPS in route finding systems, map services (Google Maps, Google Earth)              |
| <i>Science</i>                               | The Grid as an enabling technology for collaboration between scientists             |
| <i>Environmental management</i>              | sensor technology to monitor earthquakes, floods, or tsunamis                       |

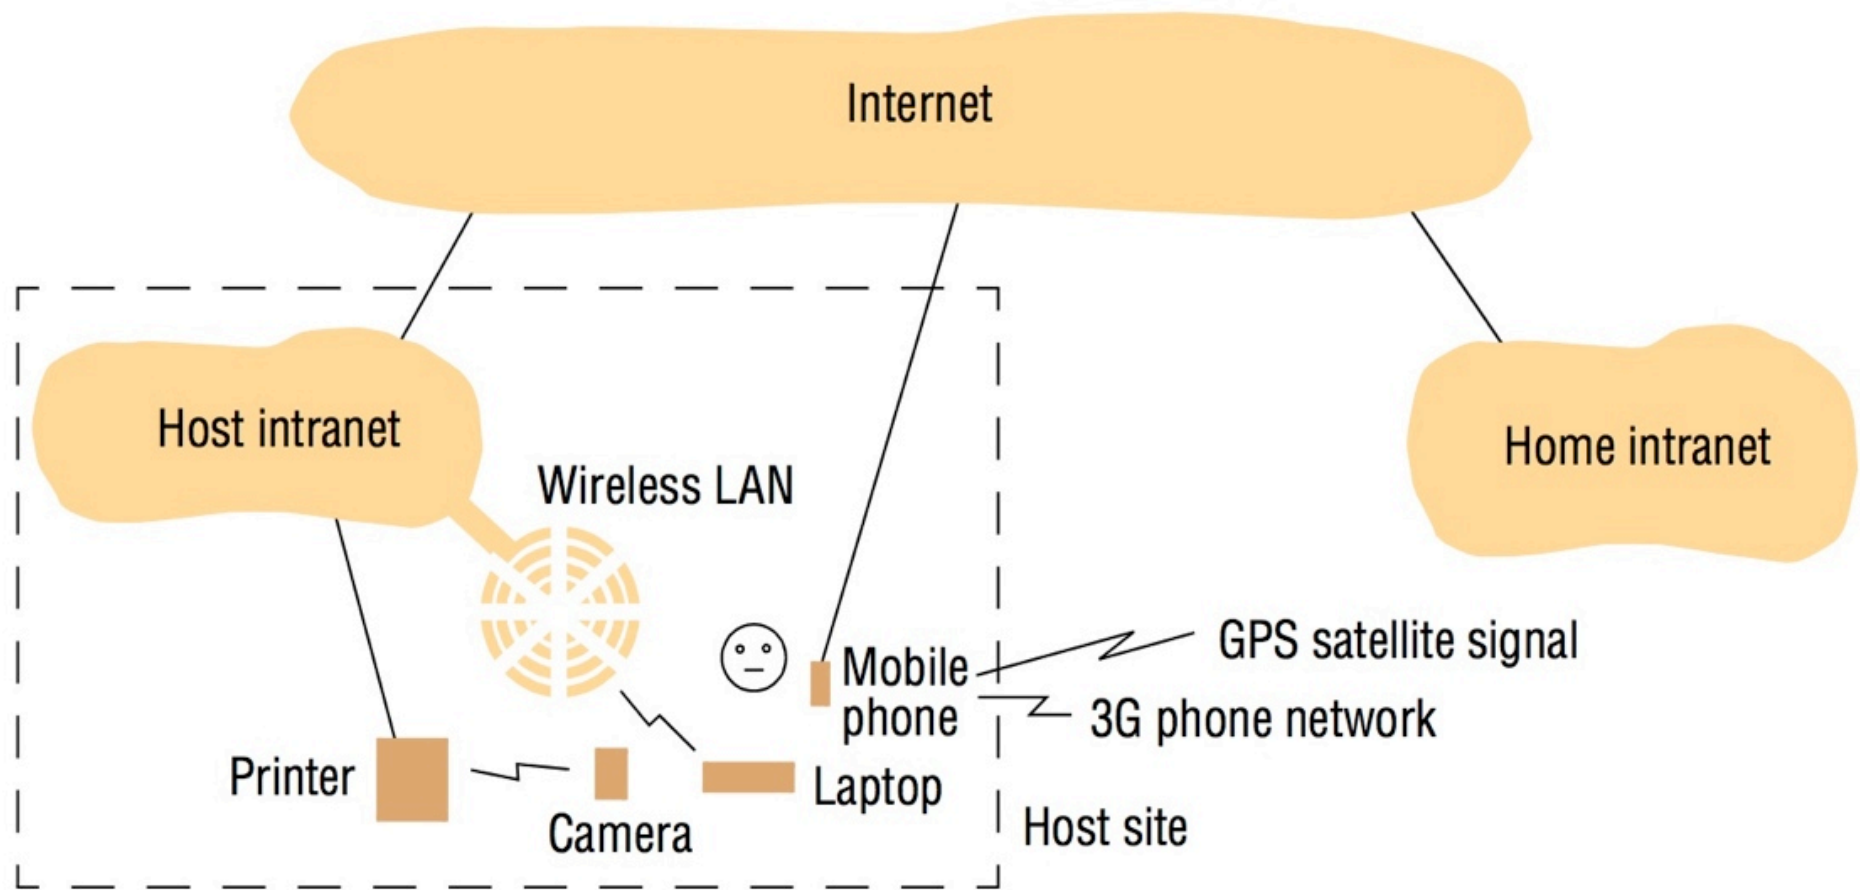
# Trends in distributed systems

- Pervasive networking and the modern Internet



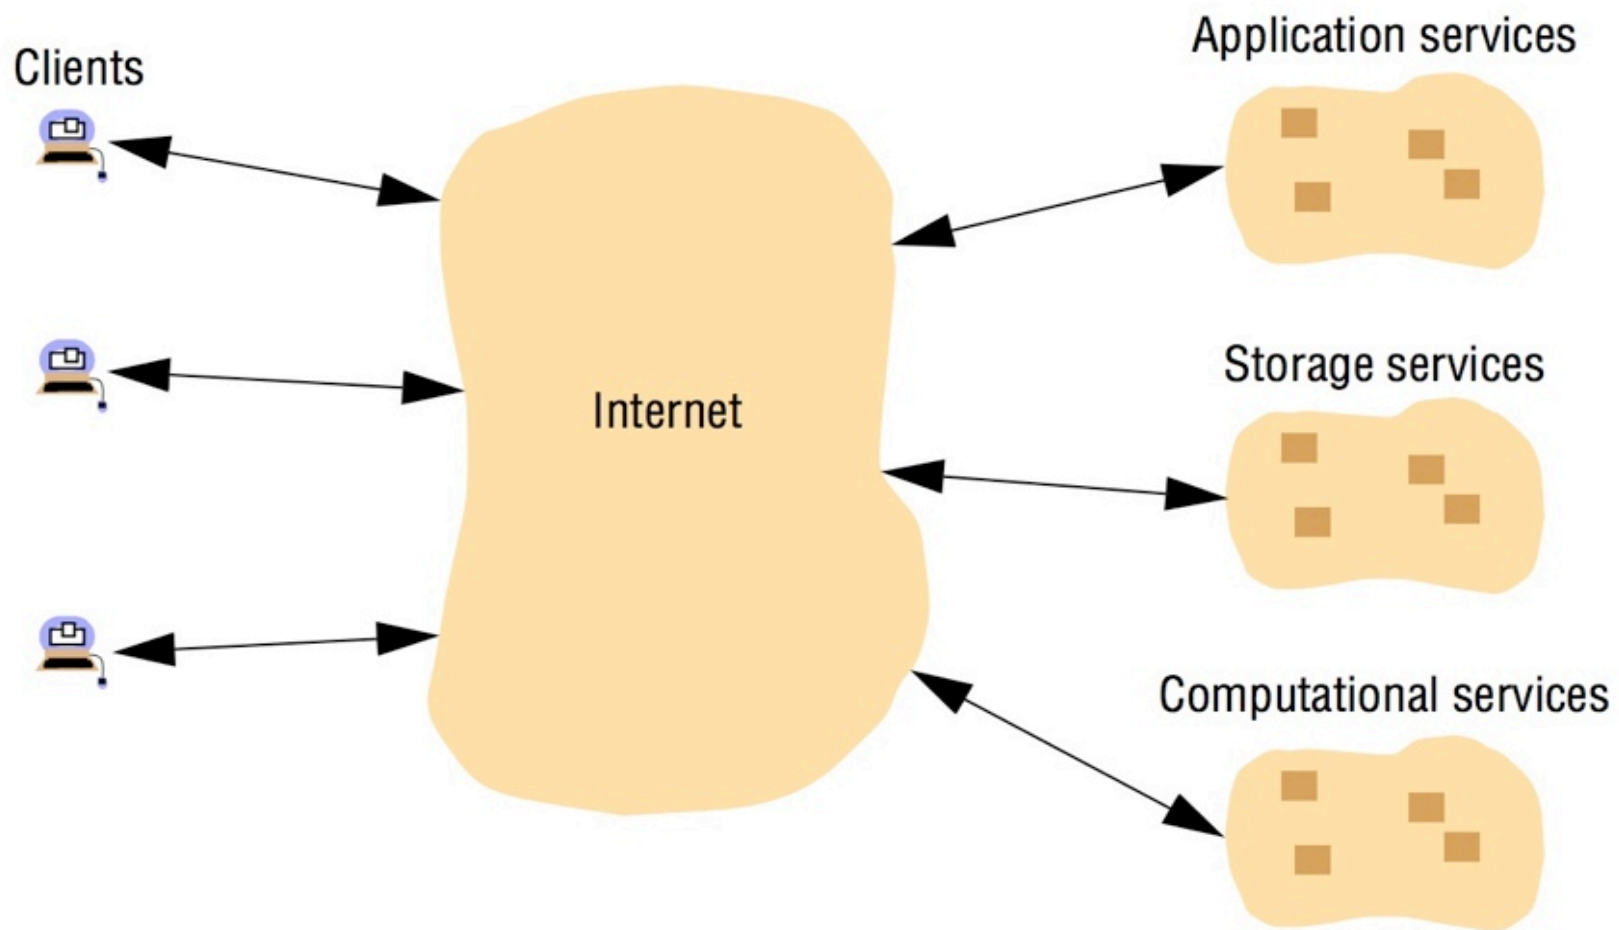
# Trends in distributed systems (cont.)

- Mobile and ubiquitous computing (Internet of Things)



# Trends in distributed systems (cont.)

- Distributed computing as a utility



# Resource sharing and the Web

- Patterns of sharing vary widely in scope and how closely users work together
  - A search engine on the Web is used by people all over the world
  - In CSCW (computer-supported cooperative working) a group of users share resources in a small, closed group

# Service, server, and client

- Requests are sent in messages from clients to a server and replies (services) are sent in messages from the server to the clients, e.g., a web browser requests a web page from a web server
- A complete interaction between a client and a server, from the point when the client sends its request to when it receives the server's response, is called a *remote invocation*. (cf. function call in an address space)
- The same process may be both a client and a server since servers sometimes invoke operations on other servers

# Building a distributed system

- “The number of people who know how to build really solid distributed systems ... is about ten.”
  - Scott Shenker, Professor at UC Berkeley
- The point: it's hard to build a solid distributed system

# Why is it hard to build one?

- **Scale**: hundreds or thousands of machines
  - Google: 4K-machine MapReduce cluster
  - Yahoo!: 4K-machine Hadoop cluster
  - Akamai: 70K machines distributed over the world
  - Facebook: 60K machines providing the service
  - Hard enough to program one machine!
- **Dynamism**: machines do fail!
  - 50 machine failures out of 20K machine cluster per day (reported by Yahoo!)
  - 1 disk failure out of 16K disks every 6 hours (reported by Google)
- **What else?**
  - Concurrent execution, consistency, etc.



# OK, but who cares?

- This is where all the actions are!
  - What are the two biggest driving forces in the computing industry for the last five years?
  - It's the cloud!
  - And smartphones!
  - And they are distributed!
- Now --- it's all about distributed systems!
  - Well...with a bit of exaggeration... ;-)

# Challenges in building distributed systems

- Heterogeneity of components
- Openness
- Security
- Scalability
- Failure handling
- Concurrency
- Transparency

# Heterogeneity

- Networks
  - Computer hardware
  - Operating systems
  - Programming languages
  - Implementations by different developers
- 
- Internet protocols
  - Middleware – CORBA, Java RMI, Apache Thrift
  - Mobile code – code sent from one computer to another and run at the destination, e.g., applets; virtual machine approach; JavaScript

# Openness

- Is the system **extensible** in various ways?
- The degree to which new resource-sharing services can be added and be made available for use by various programs
- Achieved by **publishing** the key interfaces and by conforming to a uniform communication mechanism

# Security

- Security for info resources has 3 components:
  - Confidentiality (protection against disclosure to unauthorized individuals)
  - Integrity (protection against alteration or corruption)
  - Availability (protection against interference with the means to access the resources)
- Encryption techniques are used to achieve security
- Two security challenges
  - Denial of service attacks (bombarding the service with a large number of pointless requests)
  - Security of mobile code (possible effects of running it is unpredictable)

# Scalability

- Distributed systems operate effectively and efficiently at many different scales
- A system is scalable if it will remain effective when there is a significant increase in the number of resources and users

# Growth of the Internet

| <i>Date</i> | <i>Computers</i> | <i>Web servers</i> | <i>Percentage</i> |
|-------------|------------------|--------------------|-------------------|
| 1993, July  | 1,776,000        | 130                | 0.008             |
| 1995, July  | 6,642,000        | 23,500             | 0.4               |
| 1997, July  | 19,540,000       | 1,203,096          | 6                 |
| 1999, July  | 56,218,000       | 6,598,697          | 12                |
| 2001, July  | 125,888,197      | 31,299,592         | 25                |
| 2003, July  | ~200,000,000     | 42,298,371         | 21                |
| 2005, July  | 353,284,187      | 67,571,581         | 19                |

# Challenges with scalability

- Controlling the cost of physical resources
  - For a system with  $n$  users to be scalable the quantity of resources required to support them should be at most  $O(n)$
- Controlling the performance loss
  - Managing a set of data whose size is proportional to the number of users or resources, e.g., DNS table, with a hierarchical structure thus  $O(\log n)$  performance for lookup
  - For a system to be scalable the maximum performance loss should be no worse than this
- Preventing software resources running out
  - Running out IP addresses (32 bits in 1970s vs. 128 bits being adopted requiring modifications to many software components)
- Avoiding performance bottlenecks
  - Name table in DNS was centralized in the past
  - Now partitioned between servers located throughout the Internet and administered locally



# Failure handling (next)

- Failures in a distributed system are **partial** making failure handling particularly difficult
- Detecting failures – some are difficult to detect, e.g., a remote crashed server in the Internet
- Masking failures – making detected failures less severe, e.g.,
  - Retransmit message when they fail to arrive
  - Redundant files on a pair of disks
- When one of the components fails, only the work that was using the failed component is affected

# Concurrency

- Several clients may attempt to access a shared resource at the same time, e.g., bid data in an auction
- Services and applications generally allow multiple client requests to be processed concurrently
- Any object that represents a shared resource in a distributed system must be responsible for ensuring correct semantics in a concurrent environment
- For an object to be safe in a concurrent environment its operations must be synchronized in such a way that its data remains consistent

# Transparency

- Concealing from the user and the application programmer of the separation of components in a distributed system so that the system is perceived as a whole rather than as a collection of independent components
- The ANSA Reference Manual and the International Organization for Standardization's Reference Manual for Open Distributed Processing identify eight forms of transparency

# Transparencies

- *Access transparency*: enables local and remote resources to be accessed using identical operations
- *Location transparency*: enables resources to be accessed without knowledge of their physical or network location (for example, which building or IP address)
- *Concurrency transparency*: enables several processes to operate concurrently using shared resources without interference between them
- *Replication transparency*: enables multiple instances of resources to be used to increase reliability and performance without knowledge of the replicas by users or application programmers
- *Failure transparency*: enables the concealment of faults, allowing users and application programs to complete their tasks despite the failure of hardware or software components
- *Mobility transparency*: allows the movement of resources and clients within a system without affecting the operation of users or programs
- *Performance transparency*: allows the system to be reconfigured to improve performance as loads vary
- *Scaling transparency*: allows the system and applications to expand in scale without change to the system structure or the application algorithms

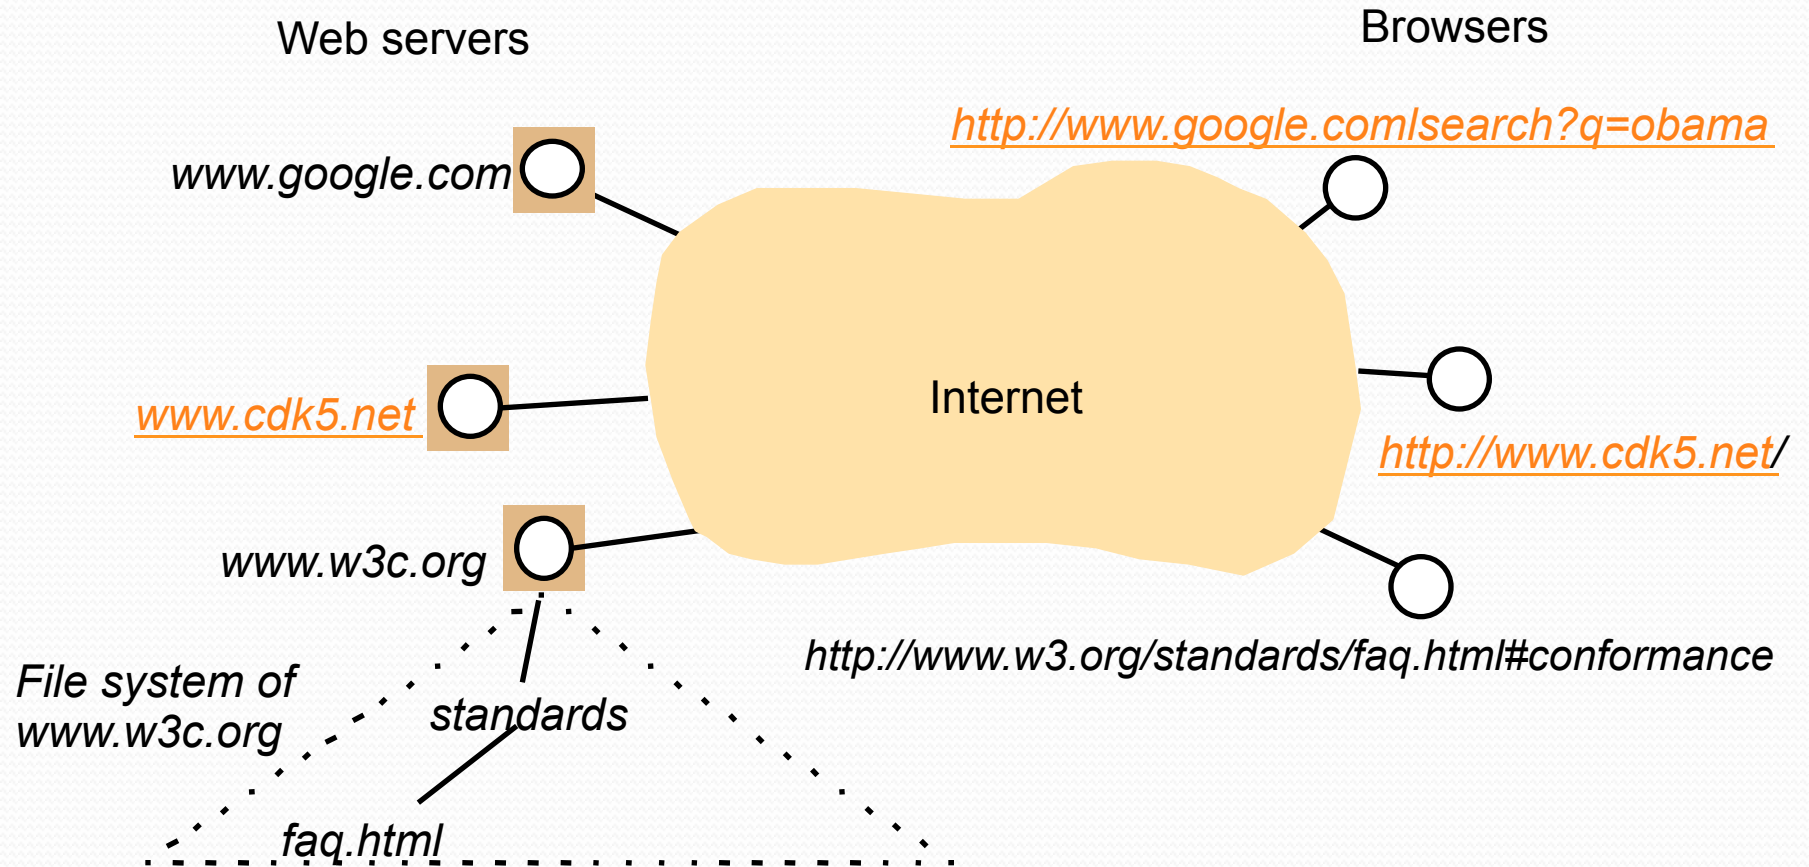
# Case study: The WWW

- The Web is an open system – can be extended in new ways without disturbing its existing functionality
  - Its operation is based on communication standards and document standards
  - Open with respect to the types of resource that can be published and shared on it

# The WWW (cont.)

- The Web is based on 3 main standard technological components:
  - HTML
  - URL's
  - A client-server architecture with HTTP

# Web servers and web browsers



# HTML (HyperText Markup Language)

- Used to specify the text and images that make up the contents of a web page
  - headings, paragraphs, tables, images, and links
- An HTML text is stored in a file that a web server can access
- A browser retrieves the file from a web server and renders the content of the file and displays
- HTML5
  - New tags: <video>, <audio>, <canvas>, <section>, <article>, <header>, <nav>, etc.
  - Eliminate plugins



# URLs (Uniform Resource Locators)

- Used to identify a resource
- The term used in web architecture documents is URI (Uniform Resource Identifier)
- The format of a URL:
  - [http://servername](http://servername[:port][pathname][?query][#fragment]) [:port][pathname][?query][#fragment]
  - For example,
    - <http://www.cmc.edu>
    - <http://www.cmc.edu/cs135/index.html#intro>
    - <http://www.google.com/search?q=lee>

# HTTP (HyperText Transfer Protocol)

- Defines the ways in which browsers and other types of clients interact with web servers
- Main features for retrieving resources
  - Request-reply interactions
  - Content types – MIME types, e.g., text/html, image/GIF
  - One resource per request
  - Simple access control (password)

# Dynamic pages

- Interacting with services that generate data rather than retrieving static data
  - Filling out a web form
  - Server has to 'process' the user's input so the client's request is a CGI program
  - Result of running the program is returned as HTML text
- Downloaded code (cf. CGI code)
  - JavaScript code
  - Applets

# Web services

- Programmatic access to web resources is commonplace, i.e., programs other than browsers can be clients of the Web too
- HTML and HTTP standards are lacking for programmatic interoperation
- XML (Extensible Markup Language) to represent standard, structured, application-specific forms
  - Meta-language for describing data – portable between applications

# Discussion of the Web

- Hugely successful
- Problems
  - Dangling links due to deleted resources
  - Users getting 'lost in hyperspace'
  - Search engines are imperfect at producing what the user specifically intends
  - The Web faces problems of scale
    - Use of caching in browsers
    - Division of the server's load across clusters of computers

# Next time

- Distributed system models